

Sai Krishna Reddy Poluri

Data Engineer

Bentonville, AR | saikrishnareddypoluri26@gmail.com | (945) 241-5664 | [LinkedIn](#) | [GitHub](#)

PROFESSIONAL SUMMARY

Databricks Certified Data Engineer with **4+ years of experience** building **scalable distributed systems**, **lakehouse architectures**, and **real-time ETL/ELT pipelines** on **Azure, GCP, and AWS**. Expertise in **Apache Spark, Databricks, and Data Modeling** across **OLAP/OLTP** systems— with a track record of **reducing cloud costs**, **accelerating data delivery**, and enabling **self-service analytics** at **enterprise scale**. Passionate about **cloud-native architectures**, **automation**, and **continuous improvement** across **data ecosystems**.

SKILLS

Big Data Processing: Apache Spark, Spark Structured Streaming, Databricks, Kafka, Hadoop, MapReduce, YARN, Hive, Sqoop

Cloud Platforms: Azure (Data Lake, Event Hubs, CosmosDB, Fabric, ADF), GCP (Dataproc, GCS, BigQuery), AWS (EC2, S3, Glue, Redshift)

Programming: Python, PySpark, SQL, Spark SQL

ETL & Orchestration: Databricks Lakeflow Jobs, Azure Data Factory, Apache Airflow, dbt

Databases & Warehousing: Snowflake, BigQuery, Synapse Analytics, Redshift, MySQL, PostgreSQL, SQL Server, MongoDB, CosmosDB (NoSQL)

DevOps & CI/CD: Git, GitHub Actions, Databricks Git Repos, Docker, Terraform

BI & Visualization: Power BI, Tableau, Looker, Databricks Genie Spaces, IBM Cognos, Microsoft Excel

EXPERIENCE

Data Engineer II | Walmart | Bentonville, AR

September 2024 – Present

- Architected **low-latency Spark Structured Streaming pipelines** on Databricks, integrating **Azure Event Hubs, Kafka, Cosmos DB, APIs**, and **Delta Lake** to process **3M+ daily transactional events**, enabling **real time analytics** for **Finance** and **Supply Chain** teams.
- Led **cloud migration** of Databricks workloads from **Microsoft Azure** to **Google Cloud Platform (ADLS → GCS, Hive Metastore → Unity Catalog)**, achieving **18% reduction in cloud operational costs** while unifying **data governance & security** across **10+ data domains**.
- Designed and implemented **CDC, SCD Type 2**, and **incremental loads** with **Medallion Architecture**, maintaining **full historical lineage** of **200M+ pricing** and **promotions records**, supporting **Operations** and **Merchandising analytics** across **InClub & Online** channels.
- Deployed **Gen AI-powered (RAG based) Genie Spaces** for **self-service analytics** over governed Delta tables, reducing **ad-hoc engineering requests** by **35%** and cutting **product/business user time-to-insight** from **hours to under 5 minutes**.
- Optimized **terabyte-scale Delta Lake models (Partitioning, Clustering)** and **Spark jobs (Cluster Tuning, Caching, AQE)**, reducing average **reporting query time** by **45%** and improving downstream **dashboard refresh rates**.
- Automated **CI/CD pipelines** via **Databricks Git Repos, Looper, and Concord**, cutting **deployment lead time** by **30%** and eliminating **manual testing & deployment errors**.

Data Analyst Intern | Redshift Technologies | Dallas, TX

June 2024 – August 2024

- Developed **SQL and Python (Pandas, NumPy, Matplotlib)** data validation scripts across **10+ business datasets**, identifying critical **data quality issues** that reduced **financial reporting errors** by **15%**.
- Built interactive **Tableau dashboards** that consolidated **KPIs** for cross-functional teams, reducing manual reporting effort by **8 hrs per week**.

Data Engineer | University of Texas at Dallas | Richardson, TX

August 2023 – May 2024

- Processed and analyzed **50GB+ transportation fleet Big Data datasets** using **PySpark** and **SQL** to identify **operational inefficiencies** in **vehicle mileage** and **fuel consumption patterns**.
- Applied **feature engineering** and **anomaly detection** on sensor data, flagging **data quality anomalies** and improving **dataset reliability** for downstream academic **machine learning research models**.
- Optimized **PySpark ETL workflows** through **partitioning** and **caching strategies**, reducing **processing runtime** by **20%** and improving **query performance**; collaborated in **Agile/Scrum** sprints and tracked deliverables using **Jira**.

Data Analyst | Cognizant | India

March 2021 – April 2022

- Analyzed **insurance claims** and **policyholder data** using **SQL** and **Python** to identify **process** and **data quality gaps**, reducing downstream **reporting errors** by **6%** across a portfolio of **500K+ policyholders**.
- Delivered **demographic** and **claims trend analysis** that contributed to a **4%** improvement in **personalized insurance product targeting**, supporting **business strategy recommendations** to stakeholders.
- Standardized **data definitions, technical documentation**, and **data quality frameworks** with **cross-functional team** collaboration, ensuring **regulatory compliance** and reducing **data inconsistencies** by **30%**, accelerating **downstream report generation**.

Business Intelligence Engineer | Manvi Enterprises | India

January 2020 – February 2021

- Developed **SQL and Python** scripts to **extract, transform, and analyze customer and sales data**, delivering **actionable retention and sales trend insights** supporting **data-driven decision making**.
- Built **Excel-based reports and dashboards**, improving **data accessibility** and reducing **manual analysis time** for **management reporting**.
- Presented **concise visual summaries of performance metrics** to **leadership**, translating **data findings** into **actionable growth strategies**.

EDUCATION

Master of Science in Data Analytics | The University of Texas at Dallas, Richardson, TX

Bachelor of Technology | SASTRA University, India

CERTIFICATIONS

- [Databricks Certified Data Engineer Associate](#)
- [AWS Certified Solutions Architect - Associate](#)
- [IBM Data Analyst Professional Certificate](#)